

Foundations of Data Science

Assignment - 2

2 - A

1. Correlation coefficients

We compute the correlation coefficient using Pearson correlation as a measure. For each feature, the correlation coefficient is computed with respect to the target attribute “Appliances”, and then stored in a list for easy access.

An important thing to note here is that even negative correlations with a high magnitude should be given proper consideration because they are strongly correlated, just in an inverse fashion. To overcome this issue, we computed the absolute valued coefficients, and then sorted the list, while keeping track of the feature that a particular coefficient represented.

Here's what the Pearson correlation measure looks like in a mathematical form:


$$r = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{(x_i - \bar{x})^2 (y_i - \bar{y})^2}}$$

Now that we have a sorted list of coefficients with their respective features, we build models where we vary the number of features taken from 1 to 26. For a model with n features, we take the first n features in the sorted coefficient list. We train each model taking the number of iterations as 10000 and setting the learning rate as 0.00001. The reason we take such a low value of the learning rate is because even after normalization, the range of the numbers in the dataset is very high valued, so if we take a larger value of the learning rate (let's suppose 0.001), during gradient descent, we will see an overshoot in the cost which will increase with the number of iterations.

The result of our experiment with the data is as follows:

	Number of features	Training Error	Testing Error
0	1.0	104.017192	98.940186
1	2.0	103.972975	98.853301
2	3.0	103.939922	98.920642
3	4.0	103.648389	98.586183
4	5.0	103.645372	98.558465
5	6.0	103.319726	98.027056
6	7.0	102.808894	97.980081
7	8.0	101.942815	97.756805
8	9.0	100.560299	96.084772
9	10.0	98.707203	95.394090
10	11.0	98.627150	95.347625
11	12.0	98.602517	95.242430
12	13.0	98.535193	95.222867
13	14.0	98.491815	95.116126
14	15.0	98.482772	95.081309
15	16.0	98.453551	95.093983
16	17.0	98.231398	94.652151
17	18.0	98.001870	94.232509
18	19.0	97.954757	94.080504
19	20.0	97.837912	94.268269
20	21.0	97.783837	94.054284
21	22.0	97.783989	94.062598
22	23.0	97.758596	94.056274
23	24.0	97.391189	93.918140
24	25.0	97.386751	93.918298
25	26.0	97.386751	93.918298

Here's a graphical representation of the same for easier understanding:



The blue line represents the training error, and the orange line represents the testing error.

Inference: One thing to note from the table is that the training and testing errors for 25 features and 26 features are identical. The reason for that is that the features 'rv1' and 'rv2' are identical. So the way that we have written the code ignores the repeated columns while building the dataset out of the most correlated features, and the errors that we see for 25 and 26 features are coming from the same model. In conclusion, we see that taking 24 features out of the entire feature set is optimal to build a regression model for the given dataset.

2. Principal Component Analysis

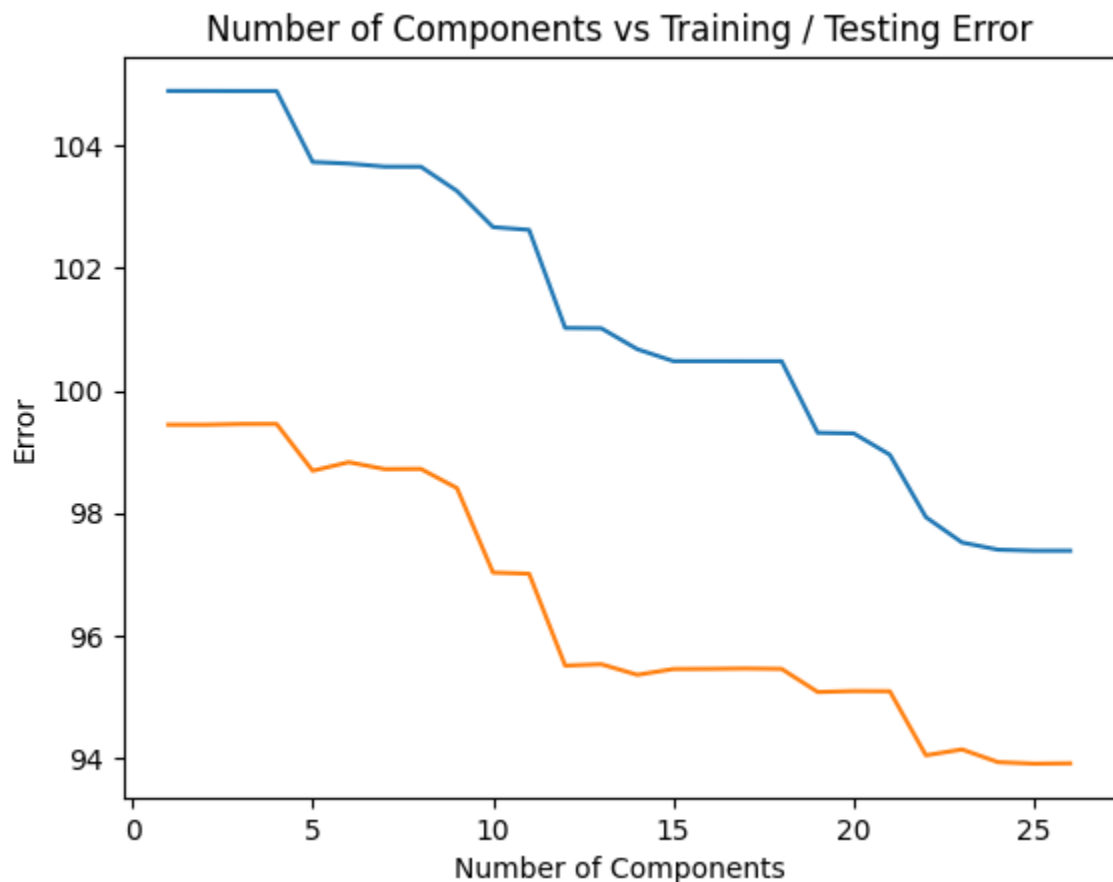
Principal Component Analysis is not a machine learning algorithm, but rather a tool that helps us to reduce the dimensions for a given feature set. In most cases, it is not possible to work with all the features from the dataset, because it becomes computationally heavy. In those kinds of situations, PCA jumps in to reduce the number of features we build our model with, by transforming our dataset into a vector space wherein we can control the number of dimensions we want to work with. We do so, by computing the eigenvalues of the dataset (only the input features), and take the eigenvectors corresponding to the first n eigenvalues in a sorted list. Now we have an n dimensional vector space wherein we can project the data points onto these n dimensions to reduce the number of features to n .

Here are the results of our experiment with the dataset:

	Number of components	Training Error	Testing Error	Percentage Variance Captured
0	1.0	104.889293	99.442722	54.873114
1	2.0	104.888877	99.442062	74.154423
2	3.0	104.887526	99.454082	81.507516
3	4.0	104.887457	99.454973	87.823240
4	5.0	103.729894	98.691059	92.539805
5	6.0	103.703555	98.831141	95.680140
6	7.0	103.653827	98.717879	97.913318
7	8.0	103.652583	98.721145	98.449016
8	9.0	103.256978	98.408648	98.811890
9	10.0	102.667292	97.030120	99.161085
10	11.0	102.621967	97.010992	99.348554
11	12.0	101.021324	95.511244	99.507935
12	13.0	101.016724	95.534441	99.633475
13	14.0	100.675335	95.360962	99.753315
14	15.0	100.480214	95.455280	99.814429
15	16.0	100.479889	95.458341	99.857296
16	17.0	100.479204	95.464531	99.890767
17	18.0	100.479050	95.458113	99.920153
18	19.0	99.311047	95.080644	99.942893
19	20.0	99.300828	95.095657	99.961881
20	21.0	98.952805	95.092514	99.975702
21	22.0	97.933973	94.047459	99.985748
22	23.0	97.521290	94.143462	99.992053
23	24.0	97.402896	93.936176	99.996645
24	25.0	97.386521	93.911110	100.000000
25	26.0	97.386077	93.916322	100.000000

We made the models by reducing the number of features to components with the number varying from 1 to 26. In addition to the training error and the testing error, we calculated the variance captured by each model. The way we computed the value of the variance was by taking the sum of the top n eigenvalues and dividing it by the sum of all eigenvalues, for a model where we consider n dimensions. The percentage of variance captured by the models vary from 54.87% to 100 %. 100 % variance being captured is intuitive because we are considering all the original features to build the model.

Here's a graphical representation of the errors for easier understanding:



The blue line represents the training error, and the orange line represents the testing error.

Inference: It turns out that the model where we have 25 features being taken into account is the one which has the lowest training and testing error.

2 - B

1. Forward feature selection

In forward feature selection, we start with an empty set of features and greedily pick the subsequent features. We are using recursion to simulate the process of picking the features. At every iteration we go through the list of features that have not been added so far, among those we pick the feature which gives us the best model using Test RMS as the metric. If during the process adding any feature does not give us a better model than the previous iteration or the previous set of features, we stop the process and return our final set of selected features.

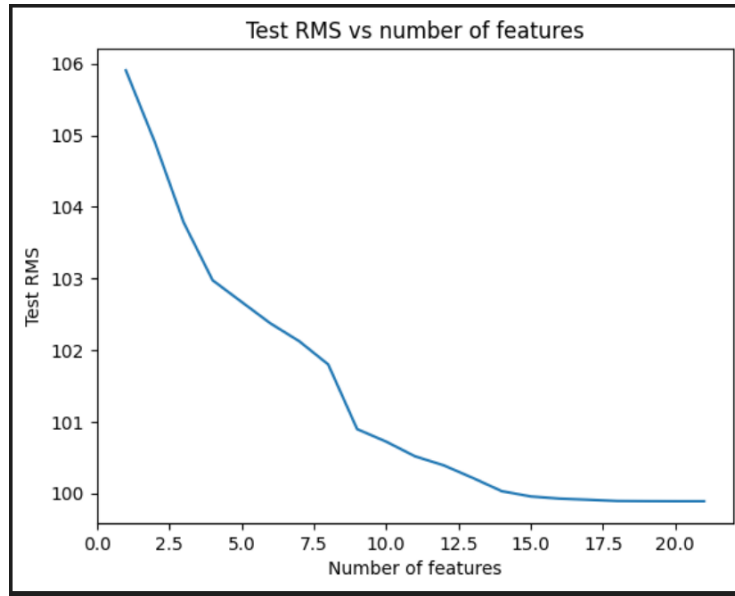
Inference: Using the greedy forward feature selection algorithm, we got the set of following 21 features as the best. Adding any new feature to the model after this point did not give us a better model as compared to the previous iteration, hence we stop at this point.

RMS : 93.64999685757

Here are the results of our experiments:

Final List of Selected Features :

```
21
['rv2', 'Tdewpoint', 'rv1', 'T4', 'RH_6', 'Visibility', 'RH_4', 'T_out', 'T2', 'T6', 'T1', 'T8',
'T9', 'T3', 'RH_3', 'RH_8', 'Windspeed', 'RH_2', 'RH_7', 'RH_1', 'RH_out']
```



No of Features selected	Train RMS	Test RMS
0	1	104.017192
1	2	103.155676
2	3	101.951968
3	4	101.449221
4	5	100.506894
5	6	100.430860
6	7	100.383972
7	8	99.563636
8	9	98.637114
9	10	98.500797
10	11	98.474310
11	12	98.348009
12	13	98.316741
13	14	98.233238
14	15	98.221206
15	16	98.207958
16	17	98.179796
17	18	98.165828
18	19	98.161910
19	20	98.151917
20	21	98.151917

2. Backward feature selection

In backward feature selection we start off with all the features given and then the features with least relevance are removed one by one. We are using recursion to simulate the process of elimination. During the first iteration we have all the features and then for each feature present we build a regression model leaving that feature out, the feature whose removal yields the best model is chosen to be eliminated. We stop the process when elimination of no process yields a better model as compared to the current model or only one feature remains.

Inference: Using the greedy backward feature selection algorithm, we got the set of following 20 features as the best. Removing any feature after this point did not give us a better model as compared to the previous iteration, hence we stop at this point.


RMS : 93.91829821626636

Here are the results of our experiments:

Final List of Selected Features :

20

```
['RH_1', 'T2', 'RH_2', 'T3', 'RH_3', 'T4', 'RH_4', 'RH_5', 'T6', 'RH_7', 'T8', 'RH_8', 'T9',  
'RH_9', 'T_out', 'Press_mm_hg', 'Windspeed', 'Visibility', 'Tdewpoint', 'rv2']
```

	No of Features selected	Train RMS	Test RMS
0	26	97.389094	93.935186
1	25	97.479512	93.790546
2	24	97.507268	93.699630
3	23	97.509033	93.660533
4	22	97.513845	93.625865
5	21	97.518902	93.614898
6	20	97.518902	93.614898

2 - C

1. Comparative Analysis of models and feature selection techniques

- The method of Correlation coefficients shows us that taking 24 features is optimal and the testing RMS for the model constructed using 24 features is 93.91
- The method of Principal Component Analysis shows that using a model which requires 25 components is optimal with a testing RMS of 93.91
- The method Forward Feature Selection shows that using a model which uses 21 of the original features to train itself returns a testing RMS of around 93.64.
- The method Backward Feature Selection shows that using a model which uses 20 of the original features to train itself returns a testing RMS of around 93.61.

Clearly, the Backward Feature Selection yields the best results out of the ones that have been computed as a part of this assignment.

Group Members :

1. Arkishman Ghosh - 2020A7PS2077H
2. Ashwin Arun - 2020A7PS1291H
3. Ashwin Naveen Pugalia - 2020A7PS1080H